

Exploring CAP Theorem, BASE, and NoSQL Databases

CLASS ACTIVITY

Ahmad Abdullah i22-1609

Abdul Sami Qasim i22-1725

1. Introduction to CAP Theorem (15 minutes)

Questions to Answer:

1. **What does the CAP theorem stand for?**

CAP refers to Consistency, Availability, and Partition Tolerance.

2. **What does each letter represent in the context of database systems?**

- **Consistency:** Ensures that all nodes in the system reflect the same data at any given moment.
- **Availability:** Guarantees that each request receives a reply, even if the request cannot be completed successfully.
- **Partition Tolerance:** Ensures the system continues functioning despite network disruptions or communication breakdowns.

3. **Why is the CAP theorem important when designing distributed systems?**

It provides a framework to understand the compromises required when balancing consistency, uptime, and fault tolerance, aiding in designing systems that meet specific needs.

4. **What is the relationship between consistency, availability, and partition tolerance?**

The theorem asserts that a distributed system can only prioritize two out of these three characteristics during a network partition.

5. **Can a system satisfy all three conditions simultaneously? Why or why not?**

It's not feasible because, during network partitioning, one must choose between providing consistent responses or remaining highly available.

6. **Examples of real-life systems:**

- **Consistency & Partition Tolerance:** HBase, MongoDB
- **Availability & Partition Tolerance:** Cassandra, DynamoDB

2. BASE Model (10 minutes)

Questions to Answer:

1. **What does BASE stand for, and how is it different from ACID?**

- **BASE:** Refers to **Basically Available**, **Soft state**, and **Eventually consistent**. It emphasizes high availability and tolerating delays in consistency.
- **ACID:** Focuses on maintaining strict **Atomicity**, **Consistency**, **Isolation**, and **Durability** for transactions.

2. **Key principles of BASE:**

- **Eventual Consistency:** Guarantees that data will eventually align across all nodes.

- **Soft State:** Allows data to change without external input due to delayed synchronization.
- **Basic Availability:** Ensures the system responds, even if the responses aren't always accurate or immediate.

3. Types of NoSQL Databases (40 minutes)

Document-based Databases:

- **Description:** Stores information as structured or semi-structured documents, typically in formats like JSON or XML.
- **Key Characteristics:**
 - Adaptable schemas, ideal for hierarchical data storage, and supports horizontal scalability.
- **Examples:** MongoDB, CouchDB
- **Use Cases:**
 - **Problem:** Managing semi-structured data.
 - **Why it fits:** Easily modifiable schemas and excellent scalability.
 - **Challenges:** Requires careful planning for queries and indexing.

Key-Value Stores:

- **Description:** Data is stored as pairs of keys and their associated values.
- **Key Characteristics:**
 - Straightforward design, optimized for rapid read and write operations, and scales efficiently.
- **Examples:** Redis, DynamoDB
- **Use Cases:**
 - **Problem:** Session management and caching.
 - **Why it fits:** Provides low-latency and high-speed lookups.
 - **Challenges:** Limited ability to handle complex queries.

Column-family Stores:

- **Description:** Data is arranged in a table-like structure, where rows can have a variable number of columns.
- **Key Characteristics:**
 - Highly scalable and suited for large-scale analytical workloads.
- **Examples:** Cassandra, HBase

- **Use Cases:**
 - **Problem:** Handling logs or time-series data.
 - **Why it fits:** Distributed nature and optimized for high write speeds.
 - **Challenges:** Maintenance and administration can be intricate.

Graph Databases:

- **Description:** Represents data as nodes and their relationships using edges.
- **Key Characteristics:**
 - Tailored for exploring and analyzing complex relationships between data points.
- **Examples:** Neo4j, Amazon Neptune
- **Use Cases:**
 - **Problem:** Social network analysis or recommendation engines.
 - **Why it fits:** Provides efficient relationship traversal.
 - **Challenges:** May not scale efficiently for extremely large datasets.

4. Critical Thinking Questions (15 minutes)

1. **How does the CAP theorem affect the choice between different NoSQL databases?**
The theorem influences system design by prioritizing either consistency, availability, or fault tolerance based on specific application needs. For instance, systems requiring real-time data access might favor availability and partition tolerance over strict consistency.
2. **Why do companies choose NoSQL over RDBMS?**
NoSQL databases offer greater flexibility with schema design, improved scalability, better performance for big data, and cost efficiency, making them ideal for modern dynamic applications.
3. **Potential risks or trade-offs of using BASE instead of ACID:**
 - **Risks:** Temporary data inconsistencies and added complexity for developers to ensure eventual consistency in application logic.
 - **Trade-offs:** While BASE sacrifices immediate consistency, it provides better scalability and system availability.